

Project1

August 5, 2022

```
[1]: # Importing the libraries
import pandas as pd
import numpy as np
import tensorflow as tf
```

```
[2]: # Importing the data
data = pd.read_csv('loan_data (1).csv')
```

```
[3]: data.shape
```

```
[3]: (307511, 122)
```

```
[4]: data.head()
```

```
[4]: SK_ID_CURR  TARGET  NAME_CONTRACT_TYPE  CODE_GENDER  FLAG_OWN_CAR  \
0      100002      1      Cash loans      M      N
1      100003      0      Cash loans      F      N
2      100004      0      Revolving loans      M      Y
3      100006      0      Cash loans      F      N
4      100007      0      Cash loans      M      N

      FLAG_OWN_REALTY  CNT_CHILDREN  AMT_INCOME_TOTAL  AMT_CREDIT  AMT_ANNUITY  \
0      Y      0      202500.0      406597.5      24700.5
1      N      0      270000.0      1293502.5      35698.5
2      Y      0      67500.0      135000.0      6750.0
3      Y      0      135000.0      312682.5      29686.5
4      Y      0      121500.0      513000.0      21865.5

...  FLAG_DOCUMENT_18  FLAG_DOCUMENT_19  FLAG_DOCUMENT_20  FLAG_DOCUMENT_21  \
0  ...      0      0      0      0
1  ...      0      0      0      0
2  ...      0      0      0      0
3  ...      0      0      0      0
4  ...      0      0      0      0

      AMT_REQ_CREDIT_BUREAU_HOUR  AMT_REQ_CREDIT_BUREAU_DAY  \
0      0.0      0.0
```

1	0.0	0.0
2	0.0	0.0
3	NaN	NaN
4	0.0	0.0

	AMT_REQ_CREDIT_BUREAU_WEEK	AMT_REQ_CREDIT_BUREAU_MON \
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	NaN	NaN
4	0.0	0.0

	AMT_REQ_CREDIT_BUREAU_QRT	AMT_REQ_CREDIT_BUREAU_YEAR
0	0.0	1.0
1	0.0	0.0
2	0.0	0.0
3	NaN	NaN
4	0.0	0.0

[5 rows x 122 columns]

```
[5]: #Check for null values in the dataset
data.isnull().sum().sort_values(ascending=False)
```

```
[5]: COMMONAREA_MEDI          214865
COMMONAREA_AVG              214865
COMMONAREA_MODE              214865
NONLIVINGAPARTMENTS_MODE    213514
NONLIVINGAPARTMENTS_MEDI    213514

...
REG_CITY_NOT_LIVE_CITY      0
LIVE_REGION_NOT_WORK_REGION 0
REG_REGION_NOT_WORK_REGION  0
HOUR_APPR_PROCESS_START     0
SK_ID_CURR                  0
Length: 122, dtype: int64
```

```
[6]: # WE can remove the columns which have more than 50% of missing values
perc = 50.0
min_count = int(((100-perc)/100)*data.shape[0] + 1)
mod_df = data.dropna( axis=1,
                      thresh=min_count)
```

```
[7]: mod_df.shape
```

```
[7]: (307511, 81)
```

```
[8]: # Imputing the missing values
```

```
[9]: mod_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 307511 entries, 0 to 307510
Data columns (total 81 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SK_ID_CURR                            307511 non-null  int64
1   TARGET                                307511 non-null  int64
2   NAME_CONTRACT_TYPE                    307511 non-null  object
3   CODE_GENDER                           307511 non-null  object
4   FLAG_OWN_CAR                           307511 non-null  object
5   FLAG_OWN_REALTY                       307511 non-null  object
6   CNT_CHILDREN                          307511 non-null  int64
7   AMT_INCOME_TOTAL                      307511 non-null  float64
8   AMT_CREDIT                            307511 non-null  float64
9   AMT_ANNUITY                           307499 non-null  float64
10  AMT_GOODS_PRICE                       307233 non-null  float64
11  NAME_TYPE_SUITE                       306219 non-null  object
12  NAME_INCOME_TYPE                     307511 non-null  object
13  NAME_EDUCATION_TYPE                  307511 non-null  object
14  NAME_FAMILY_STATUS                   307511 non-null  object
15  NAME_HOUSING_TYPE                    307511 non-null  object
16  REGION_POPULATION_RELATIVE           307511 non-null  float64
17  DAYS_BIRTH                           307511 non-null  int64
18  DAYS_EMPLOYED                        307511 non-null  int64
19  DAYS_REGISTRATION                    307511 non-null  float64
20  DAYS_ID_PUBLISH                      307511 non-null  int64
21  FLAG_MOBIL                           307511 non-null  int64
22  FLAG_EMP_PHONE                       307511 non-null  int64
23  FLAG_WORK_PHONE                      307511 non-null  int64
24  FLAG_CONT_MOBILE                     307511 non-null  int64
25  FLAG_PHONE                           307511 non-null  int64
26  FLAG_EMAIL                           307511 non-null  int64
27  OCCUPATION_TYPE                      211120 non-null  object
28  CNT_FAM_MEMBERS                      307509 non-null  float64
29  REGION_RATING_CLIENT                 307511 non-null  int64
30  REGION_RATING_CLIENT_W_CITY          307511 non-null  int64
31  WEEKDAY_APPR_PROCESS_START           307511 non-null  object
32  HOUR_APPR_PROCESS_START              307511 non-null  int64
33  REG_REGION_NOT_LIVE_REGION           307511 non-null  int64
34  REG_REGION_NOT_WORK_REGION           307511 non-null  int64
35  LIVE_REGION_NOT_WORK_REGION          307511 non-null  int64
36  REG_CITY_NOT_LIVE_CITY               307511 non-null  int64
37  REG_CITY_NOT_WORK_CITY               307511 non-null  int64
```

38	LIVE_CITY_NOT_WORK_CITY	307511	non-null	int64
39	ORGANIZATION_TYPE	307511	non-null	object
40	EXT_SOURCE_2	306851	non-null	float64
41	EXT_SOURCE_3	246546	non-null	float64
42	YEARS_BEGINEXPLUATATION_AVG	157504	non-null	float64
43	FLOORSMAX_AVG	154491	non-null	float64
44	YEARS_BEGINEXPLUATATION_MODE	157504	non-null	float64
45	FLOORSMAX_MODE	154491	non-null	float64
46	YEARS_BEGINEXPLUATATION_MEDI	157504	non-null	float64
47	FLOORSMAX_MEDI	154491	non-null	float64
48	TOTALAREA_MODE	159080	non-null	float64
49	EMERGENCYSTATE_MODE	161756	non-null	object
50	OBS_30_CNT_SOCIAL_CIRCLE	306490	non-null	float64
51	DEF_30_CNT_SOCIAL_CIRCLE	306490	non-null	float64
52	OBS_60_CNT_SOCIAL_CIRCLE	306490	non-null	float64
53	DEF_60_CNT_SOCIAL_CIRCLE	306490	non-null	float64
54	DAYS_LAST_PHONE_CHANGE	307510	non-null	float64
55	FLAG_DOCUMENT_2	307511	non-null	int64
56	FLAG_DOCUMENT_3	307511	non-null	int64
57	FLAG_DOCUMENT_4	307511	non-null	int64
58	FLAG_DOCUMENT_5	307511	non-null	int64
59	FLAG_DOCUMENT_6	307511	non-null	int64
60	FLAG_DOCUMENT_7	307511	non-null	int64
61	FLAG_DOCUMENT_8	307511	non-null	int64
62	FLAG_DOCUMENT_9	307511	non-null	int64
63	FLAG_DOCUMENT_10	307511	non-null	int64
64	FLAG_DOCUMENT_11	307511	non-null	int64
65	FLAG_DOCUMENT_12	307511	non-null	int64
66	FLAG_DOCUMENT_13	307511	non-null	int64
67	FLAG_DOCUMENT_14	307511	non-null	int64
68	FLAG_DOCUMENT_15	307511	non-null	int64
69	FLAG_DOCUMENT_16	307511	non-null	int64
70	FLAG_DOCUMENT_17	307511	non-null	int64
71	FLAG_DOCUMENT_18	307511	non-null	int64
72	FLAG_DOCUMENT_19	307511	non-null	int64
73	FLAG_DOCUMENT_20	307511	non-null	int64
74	FLAG_DOCUMENT_21	307511	non-null	int64
75	AMT_REQ_CREDIT_BUREAU_HOUR	265992	non-null	float64
76	AMT_REQ_CREDIT_BUREAU_DAY	265992	non-null	float64
77	AMT_REQ_CREDIT_BUREAU_WEEK	265992	non-null	float64
78	AMT_REQ_CREDIT_BUREAU_MON	265992	non-null	float64
79	AMT_REQ_CREDIT_BUREAU_QRT	265992	non-null	float64
80	AMT_REQ_CREDIT_BUREAU_YEAR	265992	non-null	float64

dtypes: float64(27), int64(41), object(13)

memory usage: 190.0+ MB

```
[10]: mod_df.OCCUPATION_TYPE.value_counts()
```

```
[10]: Laborers          55186
      Sales staff      32102
      Core staff       27570
      Managers         21371
      Drivers          18603
      High skill tech staff 11380
      Accountants       9813
      Medicine staff    8537
      Security staff    6721
      Cooking staff     5946
      Cleaning staff    4653
      Private service staff 2652
      Low-skill Laborers 2093
      Waiters/barmen staff 1348
      Secretaries       1305
      Realty agents     751
      HR staff          563
      IT staff          526
      Name: OCCUPATION_TYPE, dtype: int64
```

```
[11]: mod_df.OCCUPATION_TYPE.isnull().sum()
```

```
[11]: 96391
```

```
[12]: mod_df.OCCUPATION_TYPE = mod_df.OCCUPATION_TYPE.fillna('Missing')
```

/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5170:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self[name] = value

```
[13]: pd.set_option('display.float_format', lambda x: '%2f' %x)
```

```
[14]: mod_df.EXT_SOURCE_3.describe()
```

```
[14]: count    246546.000000
      mean       0.510853
      std       0.194844
      min       0.000527
      25%       0.370650
      50%       0.535276
      75%       0.669057
      max       0.896010
```

Name: EXT_SOURCE_3, dtype: float64

```
[15]: mod_df.EXT_SOURCE_2 = mod_df.EXT_SOURCE_2.fillna(mod_df.EXT_SOURCE_2.mean())
```

```
[16]: mod_df.YEARS_BEGINEXPLUATATION_AVG.describe()
```

```
[16]: count    157504.000000  
      mean      0.977735  
      std      0.059223  
      min      0.000000  
      25%      0.976700  
      50%      0.981600  
      75%      0.986600  
      max      1.000000  
      Name: YEARS_BEGINEXPLUATATION_AVG, dtype: float64
```

```
[17]: mod_df.YEARS_BEGINEXPLUATATION_AVG = mod_df.YEARS_BEGINEXPLUATATION_AVG.  
      ↪fillna(mod_df.YEARS_BEGINEXPLUATATION_AVG.mean())
```

```
[18]: mod_df.FLOORSMAX_AVG.describe()
```

```
[18]: count    154491.000000  
      mean      0.226282  
      std      0.144641  
      min      0.000000  
      25%      0.166700  
      50%      0.166700  
      75%      0.333300  
      max      1.000000  
      Name: FLOORSMAX_AVG, dtype: float64
```

```
[19]: mod_df.FLOORSMAX_AVG = mod_df.FLOORSMAX_AVG.fillna(mod_df.FLOORSMAX_AVG.mean())
```

```
[20]: mod_df.YEARS_BEGINEXPLUATATION_MODE.describe()
```

```
[20]: count    157504.000000  
      mean      0.977065  
      std      0.064575  
      min      0.000000  
      25%      0.976700  
      50%      0.981600  
      75%      0.986600  
      max      1.000000  
      Name: YEARS_BEGINEXPLUATATION_MODE, dtype: float64
```

```
[21]: mod_df.YEARS_BEGINEXPLUATATION_MODE = mod_df.YEARS_BEGINEXPLUATATION_MODE.  
      ↪fillna(mod_df.YEARS_BEGINEXPLUATATION_MODE.mean())
```

```
[22]: mod_df.FLOORSMAX_MODE.describe()
```

```
[22]: count    154491.000000  
      mean      0.222315  
      std      0.143709  
      min      0.000000  
      25%      0.166700  
      50%      0.166700  
      75%      0.333300  
      max      1.000000  
      Name: FLOORSMAX_MODE, dtype: float64
```

```
[23]: mod_df.FLOORSMAX_MODE = mod_df.FLOORSMAX_MODE.fillna(mod_df.FLOORSMAX_MODE.  
      ↪mean())
```

```
[24]: mod_df.YEARS_BEGINEXPLUATATION_MEDI.describe()
```

```
[24]: count    157504.000000  
      mean      0.977752  
      std      0.059897  
      min      0.000000  
      25%      0.976700  
      50%      0.981600  
      75%      0.986600  
      max      1.000000  
      Name: YEARS_BEGINEXPLUATATION_MEDI, dtype: float64
```

```
[25]: mod_df.YEARS_BEGINEXPLUATATION_MEDI = mod_df.YEARS_BEGINEXPLUATATION_MEDI.  
      ↪fillna(mod_df.YEARS_BEGINEXPLUATATION_MEDI.median())
```

```
[26]: mod_df.FLOORSMAX_MEDI.describe()
```

```
[26]: count    154491.000000  
      mean      0.225897  
      std      0.145067  
      min      0.000000  
      25%      0.166700  
      50%      0.166700  
      75%      0.333300  
      max      1.000000  
      Name: FLOORSMAX_MEDI, dtype: float64
```

```
[27]: mod_df.FLOORSMAX_MEDI = mod_df.FLOORSMAX_MEDI.fillna(mod_df.FLOORSMAX_MEDI.  
      ↪median())
```

```
[28]: mod_df.TOTALAREA_MODE.describe()
```

```
[28]: count    159080.000000
      mean      0.102547
      std       0.107462
      min       0.000000
      25%       0.041200
      50%       0.068800
      75%       0.127600
      max       1.000000
      Name: TOTALAREA_MODE, dtype: float64
```

```
[29]: mod_df.TOTALAREA_MODE = mod_df.TOTALAREA_MODE.fillna(mod_df.TOTALAREA_MODE.
      ↪median())
```

```
[30]: mod_df.EMERGENCYSTATE_MODE.value_counts(dropna=False)
```

```
[30]: No      159428
      NaN     145755
      Yes      2328
      Name: EMERGENCYSTATE_MODE, dtype: int64
```

```
[31]: mod_df.EMERGENCYSTATE_MODE = mod_df.EMERGENCYSTATE_MODE.fillna('Not known')
```

```
[32]: mod_df.AMT_REQ_CREDIT_BUREAU_HOUR.describe()
```

```
[32]: count    265992.000000
      mean      0.006402
      std       0.083849
      min       0.000000
      25%       0.000000
      50%       0.000000
      75%       0.000000
      max       4.000000
      Name: AMT_REQ_CREDIT_BUREAU_HOUR, dtype: float64
```

```
[33]: mod_df.AMT_REQ_CREDIT_BUREAU_HOUR.value_counts(dropna=False)
```

```
[33]: 0.000000    264366
      nan       41519
      1.000000    1560
      2.000000     56
      3.000000     9
      4.000000     1
      Name: AMT_REQ_CREDIT_BUREAU_HOUR, dtype: int64
```

```
[34]: mod_df.AMT_REQ_CREDIT_BUREAU_HOUR = mod_df.AMT_REQ_CREDIT_BUREAU_HOUR.fillna(0.
      ↪0)
```



```
[35]: mod_df.AMT_REQ_CREDIT_BUREAU_DAY.value_counts(dropna=False)
```

```
[35]: 0.000000    264503
      nan      41519
      1.000000    1292
      2.000000     106
      3.000000     45
      4.000000     26
      5.000000      9
      6.000000      8
      9.000000      2
      8.000000      1
      Name: AMT_REQ_CREDIT_BUREAU_DAY, dtype: int64
```

```
[36]: mod_df.AMT_REQ_CREDIT_BUREAU_DAY = mod_df.AMT_REQ_CREDIT_BUREAU_DAY.fillna(0.0)
```

```
[37]: mod_df.AMT_REQ_CREDIT_BUREAU_WEEK.value_counts(dropna=False)
```

```
[37]: 0.000000    257456
      nan      41519
      1.000000    8208
      2.000000     199
      3.000000      58
      4.000000      34
      6.000000      20
      5.000000      10
      8.000000       5
      7.000000       2
      Name: AMT_REQ_CREDIT_BUREAU_WEEK, dtype: int64
```

```
[38]: mod_df.AMT_REQ_CREDIT_BUREAU_WEEK = mod_df.AMT_REQ_CREDIT_BUREAU_WEEK.fillna(0.
      ↪0)
```

```
[39]: mod_df.AMT_REQ_CREDIT_BUREAU_MON.value_counts(dropna=False)
```

```
[39]: 0.000000    222233
      nan      41519
      1.000000    33147
      2.000000    5386
      3.000000    1991
      4.000000    1076
      5.000000     602
      6.000000     343
      7.000000     298
      9.000000     206
      8.000000     185
      10.000000    132
```

11.000000	119
12.000000	77
13.000000	72
14.000000	40
15.000000	35
16.000000	23
17.000000	14
18.000000	6
19.000000	3
23.000000	1
27.000000	1
22.000000	1
24.000000	1

Name: AMT_REQ_CREDIT_BUREAU_MON, dtype: int64

```
[40]: mod_df.AMT_REQ_CREDIT_BUREAU_MON = mod_df.AMT_REQ_CREDIT_BUREAU_MON.fillna(0.0)
```

```
[41]: mod_df.AMT_REQ_CREDIT_BUREAU_QRT.value_counts(dropna=False)
```

```
[41]: 0.000000    215417
      nan        41519
      1.000000    33862
      2.000000    14412
      3.000000     1717
      4.000000     476
      5.000000      64
      6.000000     28
      7.000000      7
      8.000000      7
      19.000000      1
      261.000000      1
      Name: AMT_REQ_CREDIT_BUREAU_QRT, dtype: int64
```

```
[42]: mod_df.AMT_REQ_CREDIT_BUREAU_QRT = mod_df.AMT_REQ_CREDIT_BUREAU_QRT.fillna(0.0)
```

```
[43]: mod_df.AMT_REQ_CREDIT_BUREAU_YEAR.value_counts(dropna=False)
```

```
[43]: 0.000000    71801
      1.000000    63405
      2.000000    50192
      nan        41519
      3.000000    33628
      4.000000    20714
      5.000000    12052
      6.000000     6967
      7.000000     3869
      8.000000     2127
```

```

9.000000    1096
11.000000     31
12.000000     30
10.000000     22
13.000000     19
14.000000     10
17.000000      7
15.000000      6
19.000000      4
18.000000      4
16.000000      3
21.000000      1
23.000000      1
25.000000      1
20.000000      1
22.000000      1
Name: AMT_REQ_CREDIT_BUREAU_YEAR, dtype: int64

```

```
[44]: mod_df.AMT_REQ_CREDIT_BUREAU_YEAR = mod_df.AMT_REQ_CREDIT_BUREAU_YEAR.fillna(0.
      ↪0)
```

```
[45]: mod_df.isnull().sum().sort_values(ascending=False)
```

```

[45]: EXT_SOURCE_3          60965
      NAME_TYPE_SUITE       1292
      DEF_60_CNT_SOCIAL_CIRCLE  1021
      OBS_60_CNT_SOCIAL_CIRCLE  1021
      DEF_30_CNT_SOCIAL_CIRCLE  1021

      ...
      FLAG_DOCUMENT_3         0
      FLAG_DOCUMENT_4         0
      FLAG_DOCUMENT_5         0
      FLAG_DOCUMENT_6         0
      SK_ID_CURR              0
      Length: 81, dtype: int64

```

```
[46]: mod_df = mod_df.dropna(axis=0)
```

```
[47]: mod_df.shape
```

```
[47]: (244708, 81)
```

```
[48]: # we still have enough data to train our model
```

```
[49]: #Print percentage of default to payer of the dataset for the TARGET column
```

```
[50]: default_to_payer = (mod_df.TARGET.value_counts()[1]) / (mod_df.TARGET.  
    ↪value_counts()[0]) * 100  
default_to_payer
```

```
[50]: 8.450148687516897
```

```
[51]: #default to payer percentage is 8.45%  
# here we can see that the data is highly imbalanced.
```


```
[52]: import seaborn as sns
```

```
[53]: sns.countplot(mod_df['TARGET'])
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.

FutureWarning

```
[53]: <AxesSubplot:xlabel='TARGET', ylabel='count'>
```



```
[54]: # Before we treat the imbalance in our data we need to split the data into  
    ↪train and test set so that the originality of testing set is not compromised  
# Before we split the data we need to label encode the categorical columns so  
    ↪that we will have the same columns in train and test dataset.
```

```

[55]: X = mod_df.drop('TARGET',axis =1)

[56]: y = mod_df.TARGET

[57]: X.dtypes.value_counts()

[57]: int64      40
      float64   27
      object    13
      dtype: int64

[58]: # Label encode the categorical columns

[59]: cat_cols = X.select_dtypes(include='object')

[60]: cat_cols_encoded = pd.get_dummies(cat_cols,prefix_sep='_')

[61]: cat_cols_encoded.shape

[61]: (244708, 126)

[62]: num_cols = X.select_dtypes(exclude='object')

[63]: X_encoded = pd.concat([cat_cols_encoded,num_cols],axis=1)

[64]: X_encoded.dtypes.value_counts()

[64]: uint8      126
      int64     40
      float64   27
      dtype: int64

[65]: # Splitting the model
      from sklearn.model_selection import train_test_split

[66]: X_train,X_test,y_train,y_test = train_test_split(X_encoded,y,test_size=0.
      ↪2,random_state=22)

[67]: # Now we will combine X train and y train and over sample the data to handle
      ↪the imbalance

[68]: imbalanced_data = pd.concat([X_train,y_train],axis=1)

[69]: # Split into majority and minority data
      df_majority = imbalanced_data[imbalanced_data.TARGET==0]
      df_minority = imbalanced_data[imbalanced_data.TARGET==1]

```

```
[70]: # Upsample minority class
      from sklearn.utils import resample
```

```
[71]: df_upsampled_minority =
      ↪resample(df_minority,replace=True,n_samples=180596,random_state=123)
```

```
[72]: df_upsampled_minority.shape
```

```
[72]: (180596, 194)
```

```
[73]: df_majority.shape
```

```
[73]: (180596, 194)
```

```
[74]: df_upsampled = pd.concat([df_majority, df_upsampled_minority])
```

```
[75]: df_upsampled.shape
```

```
[75]: (361192, 194)
```


```
[76]: df_upsampled.TARGET.value_counts()
```

```
[76]: 1    180596
      0    180596
      Name: TARGET, dtype: int64
```

```
[77]: sns.countplot(df_upsampled['TARGET'])
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
FutureWarning

```
[77]: <AxesSubplot:xlabel='TARGET', ylabel='count'>
```



```
[78]: # Now our data is balanced
```

```
[79]: # Getting the data ready to be fit in model
```

```
[80]: df_upsampled_X = df_upsampled.drop(['TARGET', 'SK_ID_CURR'], axis=1)
```

```
[81]: df_upsampled_y = df_upsampled.TARGET
```

```
[82]: X_test = X_test.drop('SK_ID_CURR', axis=1)
```

```
[83]: df_upsampled_X.shape
```

```
[83]: (361192, 192)
```

```
[84]: #Create model
```

```
[85]: model = tf.keras.models.Sequential()
```

```
[86]: model.add(tf.keras.layers.Reshape((192,), input_shape=(192,)))
```

```
[87]: model.add(tf.keras.layers.BatchNormalization())
```

```
[88]: model.add(tf.keras.layers.Dense(200, activation='relu'))  
model.add(tf.keras.layers.BatchNormalization())
```

```
[89]: model.add(tf.keras.layers.Dense(100, activation='relu'))
      model.add(tf.keras.layers.BatchNormalization())

[90]: model.add(tf.keras.layers.Dense(60, activation='relu'))
      model.add(tf.keras.layers.BatchNormalization())

[91]: model.add(tf.keras.layers.Dense(30, activation='relu'))
      model.add(tf.keras.layers.BatchNormalization())

[92]: #Output layer
      model.add(tf.keras.layers.Dense(1, activation='sigmoid'))

[93]: sgd_optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
      model.compile(optimizer=sgd_optimizer, loss='binary_crossentropy',
      ↪metrics=['accuracy'])

[94]: model.
      ↪fit(df_upsampled_X,df_upsampled_y,validation_data=(X_test,y_test),epochs=100,batch_size=32)
```

```
Epoch 1/100
11288/11288 [=====] - 29s 2ms/step - loss: 0.6000 -
accuracy: 0.6794 - val_loss: 0.5933 - val_accuracy: 0.6813
Epoch 2/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.5657 -
accuracy: 0.7079 - val_loss: 0.5886 - val_accuracy: 0.6881
Epoch 3/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.5386 -
accuracy: 0.7285 - val_loss: 0.5921 - val_accuracy: 0.6893
Epoch 4/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.5148 -
accuracy: 0.7458 - val_loss: 0.5850 - val_accuracy: 0.6948
Epoch 5/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.4944 -
accuracy: 0.7596 - val_loss: 0.5787 - val_accuracy: 0.7035
Epoch 6/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.4783 -
accuracy: 0.7707 - val_loss: 0.6032 - val_accuracy: 0.6944
Epoch 7/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.4620 -
accuracy: 0.7820 - val_loss: 0.5817 - val_accuracy: 0.7075
Epoch 8/100
11288/11288 [=====] - 29s 3ms/step - loss: 0.4508 -
accuracy: 0.7887 - val_loss: 0.5777 - val_accuracy: 0.7197
Epoch 9/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.4401 -
accuracy: 0.7953 - val_loss: 0.5729 - val_accuracy: 0.7273
Epoch 10/100
```


11288/11288 [=====] - 28s 2ms/step - loss: 0.4301 -
accuracy: 0.8019 - val_loss: 0.5973 - val_accuracy: 0.7191
Epoch 11/100
11288/11288 [=====] - 30s 3ms/step - loss: 0.4212 -
accuracy: 0.8069 - val_loss: 0.5805 - val_accuracy: 0.7319
Epoch 12/100
11288/11288 [=====] - 29s 3ms/step - loss: 0.4139 -
accuracy: 0.8113 - val_loss: 0.5801 - val_accuracy: 0.7318
Epoch 13/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.4070 -
accuracy: 0.8155 - val_loss: 0.5795 - val_accuracy: 0.7345
Epoch 14/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.3996 -
accuracy: 0.8193 - val_loss: 0.5828 - val_accuracy: 0.7341
Epoch 15/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.3927 -
accuracy: 0.8229 - val_loss: 0.5742 - val_accuracy: 0.7404
Epoch 16/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.3877 -
accuracy: 0.8260 - val_loss: 0.5847 - val_accuracy: 0.7386
Epoch 17/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.3832 -
accuracy: 0.8296 - val_loss: 0.5750 - val_accuracy: 0.7496
Epoch 18/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3776 -
accuracy: 0.8317 - val_loss: 0.5845 - val_accuracy: 0.7438
Epoch 19/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3747 -
accuracy: 0.8335 - val_loss: 0.5807 - val_accuracy: 0.7459
Epoch 20/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3688 -
accuracy: 0.8369 - val_loss: 0.5765 - val_accuracy: 0.7520
Epoch 21/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3655 -
accuracy: 0.8381 - val_loss: 0.5843 - val_accuracy: 0.7567
Epoch 22/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3622 -
accuracy: 0.8397 - val_loss: 0.5862 - val_accuracy: 0.7516
Epoch 23/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3592 -
accuracy: 0.8424 - val_loss: 0.5825 - val_accuracy: 0.7557
Epoch 24/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3552 -
accuracy: 0.8444 - val_loss: 0.5889 - val_accuracy: 0.7539
Epoch 25/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.3520 -
accuracy: 0.8458 - val_loss: 0.5891 - val_accuracy: 0.7596
Epoch 26/100

11288/11288 [=====] - 27s 2ms/step - loss: 0.3498 -
accuracy: 0.8468 - val_loss: 0.5870 - val_accuracy: 0.7548
Epoch 27/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3476 -
accuracy: 0.8476 - val_loss: 0.5966 - val_accuracy: 0.7581
Epoch 28/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3442 -
accuracy: 0.8500 - val_loss: 0.5915 - val_accuracy: 0.7605
Epoch 29/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.3423 -
accuracy: 0.8511 - val_loss: 0.5907 - val_accuracy: 0.7596
Epoch 30/100
11288/11288 [=====] - 28s 2ms/step - loss: 0.3392 -
accuracy: 0.8530 - val_loss: 0.5852 - val_accuracy: 0.7594
Epoch 31/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3385 -
accuracy: 0.8533 - val_loss: 0.5996 - val_accuracy: 0.7578
Epoch 32/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3356 -
accuracy: 0.8547 - val_loss: 0.5933 - val_accuracy: 0.7611
Epoch 33/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3277 -
accuracy: 0.8590 - val_loss: 0.5918 - val_accuracy: 0.7670
Epoch 38/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3249 -
accuracy: 0.8602 - val_loss: 0.5962 - val_accuracy: 0.7704
Epoch 39/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3226 -
accuracy: 0.8615 - val_loss: 0.5748 - val_accuracy: 0.7739
Epoch 40/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3212 -
accuracy: 0.8621 - val_loss: 0.5971 - val_accuracy: 0.7693
Epoch 41/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3192 -
accuracy: 0.8633 - val_loss: 0.5939 - val_accuracy: 0.7676
Epoch 42/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3195 -
accuracy: 0.8631 - val_loss: 0.6066 - val_accuracy: 0.7614
Epoch 43/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3177 -
accuracy: 0.8635 - val_loss: 0.6094 - val_accuracy: 0.7631
Epoch 44/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3153 -
accuracy: 0.8649 - val_loss: 0.5927 - val_accuracy: 0.7728
Epoch 45/100
11288/11288 [=====] - 27s 2ms/step - loss: 0.3157 -
accuracy: 0.8647 - val_loss: 0.5890 - val_accuracy: 0.7718
Epoch 46/100

11288/11288 [=====] - 28s 3ms/step - loss: 0.3151 -
 accuracy: 0.8648 - val_loss: 0.6027 - val_accuracy: 0.7698
 Epoch 47/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.3132 -
 accuracy: 0.8659 - val_loss: 0.6095 - val_accuracy: 0.7666
 Epoch 48/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.3108 -
 accuracy: 0.8674 - val_loss: 0.5981 - val_accuracy: 0.7691
 Epoch 49/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.3102 -
 accuracy: 0.8679 - val_loss: 0.5814 - val_accuracy: 0.7753
 Epoch 50/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.3095 -
 accuracy: 0.8684 - val_loss: 0.6002 - val_accuracy: 0.7680
 Epoch 51/100
 11288/11288 [=====] - 28s 2ms/step - loss: 0.3093 -
 accuracy: 0.8679 - val_loss: 0.5952 - val_accuracy: 0.7763
 Epoch 52/100
 11288/11288 [=====] - 28s 2ms/step - loss: 0.3070 -
 accuracy: 0.8688 - val_loss: 0.5963 - val_accuracy: 0.7752
 Epoch 53/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.3057 -
 accuracy: 0.8699 - val_loss: 0.5998 - val_accuracy: 0.7778
 Epoch 54/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2886 -
 accuracy: 0.8785 - val_loss: 0.6027 - val_accuracy: 0.7831
 Epoch 78/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2878 -
 accuracy: 0.8789 - val_loss: 0.6154 - val_accuracy: 0.7811
 Epoch 79/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2874 -
 accuracy: 0.8790 - val_loss: 0.6008 - val_accuracy: 0.7826
 Epoch 80/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2876 -
 accuracy: 0.8783 - val_loss: 0.6065 - val_accuracy: 0.7776
 Epoch 81/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2873 -
 accuracy: 0.8787 - val_loss: 0.6258 - val_accuracy: 0.7722
 Epoch 82/100
 11288/11288 [=====] - 28s 2ms/step - loss: 0.2863 -
 accuracy: 0.8797 - val_loss: 0.6061 - val_accuracy: 0.7823
 Epoch 83/100
 11288/11288 [=====] - 28s 2ms/step - loss: 0.2845 -
 accuracy: 0.8804 - val_loss: 0.6021 - val_accuracy: 0.7837
 Epoch 84/100
 11288/11288 [=====] - 28s 2ms/step - loss: 0.2850 -
 accuracy: 0.8797 - val_loss: 0.6147 - val_accuracy: 0.7819
 Epoch 85/100

11288/11288 [=====] - 28s 2ms/step - loss: 0.2843 -
 accuracy: 0.8801 - val_loss: 0.6031 - val_accuracy: 0.7809
 Epoch 86/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2829 -
 accuracy: 0.8811 - val_loss: 0.5967 - val_accuracy: 0.7857
 Epoch 87/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2835 -
 accuracy: 0.8808 - val_loss: 0.5904 - val_accuracy: 0.7892
 Epoch 88/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2834 -
 accuracy: 0.8808 - val_loss: 0.5915 - val_accuracy: 0.7885
 Epoch 89/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2824 -
 accuracy: 0.8812 - val_loss: 0.5991 - val_accuracy: 0.7811
 Epoch 90/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2808 -
 accuracy: 0.8821 - val_loss: 0.5964 - val_accuracy: 0.7867
 Epoch 91/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2805 -
 accuracy: 0.8823 - val_loss: 0.6179 - val_accuracy: 0.7784
 Epoch 92/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2816 -
 accuracy: 0.8815 - val_loss: 0.5890 - val_accuracy: 0.7879
 Epoch 93/100
 11288/11288 [=====] - 28s 2ms/step - loss: 0.2810 -
 accuracy: 0.8819 - val_loss: 0.5924 - val_accuracy: 0.7869
 Epoch 94/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2802 -
 accuracy: 0.8823 - val_loss: 0.6075 - val_accuracy: 0.7849
 Epoch 95/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2797 -
 accuracy: 0.8822 - val_loss: 0.6003 - val_accuracy: 0.7840
 Epoch 96/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2799 -
 accuracy: 0.8827 - val_loss: 0.5950 - val_accuracy: 0.7892
 Epoch 97/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2789 -
 accuracy: 0.8837 - val_loss: 0.5875 - val_accuracy: 0.7894
 Epoch 98/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2794 -
 accuracy: 0.8823 - val_loss: 0.5968 - val_accuracy: 0.7847
 Epoch 99/100
 11288/11288 [=====] - 28s 2ms/step - loss: 0.2779 -
 accuracy: 0.8834 - val_loss: 0.5967 - val_accuracy: 0.7919
 Epoch 100/100
 11288/11288 [=====] - 27s 2ms/step - loss: 0.2784 -
 accuracy: 0.8838 - val_loss: 0.6052 - val_accuracy: 0.7861

```

[94]: <keras.callbacks.History at 0x7fcb17193810>

[95]: y_pred = model.predict(X_test)

[96]: # Calculate Sensitivity as a metrice

[97]: m = tf.keras.metrics.Recall()

[98]: m.update_state(y_test,y_pred)

[99]: m.result().numpy()

[99]: 0.36438286

[100]: # Sensitivity of the model is 0.36

[101]: #Calculate area under receiver operating characteristics curve

[103]: m1 = tf.keras.metrics.AUC(num_thresholds=200,curve="ROC")

[104]: m1.update_state(y_test,y_pred)

[105]: m1.result().numpy()

[105]: 0.6381495

[106]: # The AUC-ROC is 0.63


[107]: from sklearn.metrics import roc_curve
      y_pred_keras = y_pred.ravel()
      fpr_keras, tpr_keras, thresholds_keras = roc_curve(y_test, y_pred_keras)

[108]: from sklearn.metrics import auc
      auc_keras = auc(fpr_keras, tpr_keras)

[109]: import matplotlib.pyplot as plt
      %matplotlib inline

[110]: plt.plot(fpr_keras, tpr_keras, label='Keras (area = {:.3f})'.format(auc_keras))
      plt.xlabel('False positive rate')
      plt.ylabel('True positive rate')
      plt.title('ROC curve')
      plt.legend(loc='best')
      plt.show()

```



[]: